

szczecin::cpp

14.03.2019

Say hello to: C++ Coroutines

PRESENTER:

JACEK NIJAKI <jacek.nijaki@siili.com>

What we will talk about?

Coroutines-TS (Technical Specification - N4775)

Included in C++20 draft (Kona, Feb 2019)
by Gor Nishanov <gorn@microsoft.com>

```
std::experimental  
#include <experimental/coroutine>
```

What is a Coroutine?



Generalization of a function

Function:

- can be invoked
- can return to the caller

Coroutine, can the same, and more:

- can suspend execution, and return
- later can be resumed from the point it was previously suspended



How does it work? - activation frame


Coroutine can be suspended w/o destroying – stack can't be used

Coroutine frame is stored on heap

Optimization possible – if lifetime nested in the frame of caller

Activation frame = stack frame + coroutine (heap) frame


Values are copied or moved to the heap frame



How does it work? - actions

Invoke


- From caller perspective no difference from function
- Returns to caller when suspended or completed
- Stack frame created same as during ordinary function call
- Finally coroutine creates frame on heap
- Copies params, return address etc. to the coroutine frame



How does it work? - actions

Suspend


- Possible on defined suspension points
- Drops state to the coroutine frame (including suspension point address)
- When prepared for resumption – considered ‘suspended’
- Place for additional logic before returning to the caller (e.g. start real I/O operation)
- Stack frame part is then freed
- Coroutine handle is transferred to the caller (can be used for resumption)



How does it work? - actions

Resume


- Done by calling `resume()` on coroutine handle
- Effectively calls into middle of the function
- Stores stack-frame with caller return address (like normal function)
- Instead of starting the function, execution passed to the point stored in the coroutine-frame



How does it work? - actions

Destroy


- Destroys coroutine frame w/o resuming
- Can be only performed on suspended coroutine
- Re-activates coroutine activation frame,
- No execution transfer, instead destructors called for variables in scope of last suspension point



How does it work? - actions

Return

- Different than normal
- Return value is stored somewhere (coroutine specific)
- Can execute additional logic before transferring execution
- Then performs **Suspend** (keeps frame alive) or **Destroy**
- Execution transferred to the caller
- Value passed to the return operation \neq return-value from the coroutine call



SAMPLE: function **fun()**, calls a coroutine **coro(int param)**

1. before



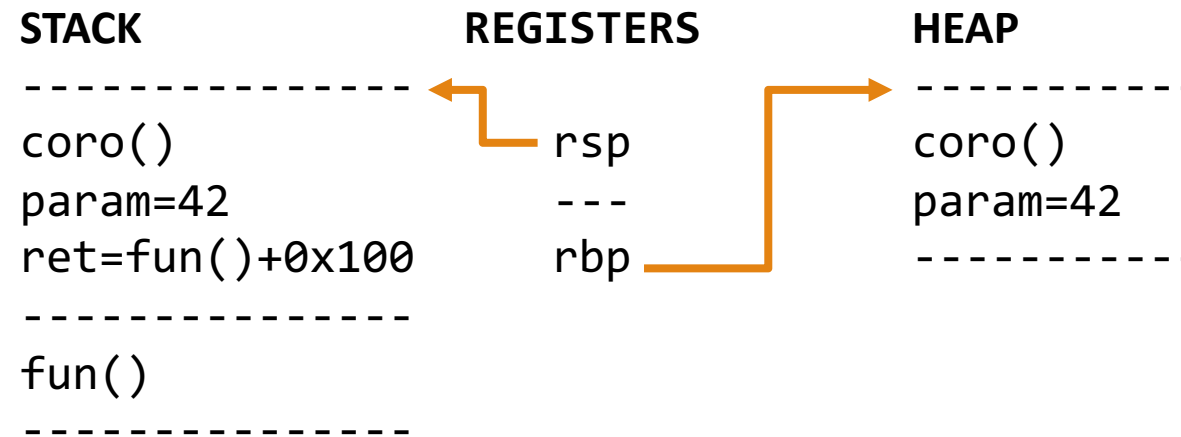
SAMPLE: function **fun()**, calls a coroutine **coro(int param)**

2. **coro(42)** is called:



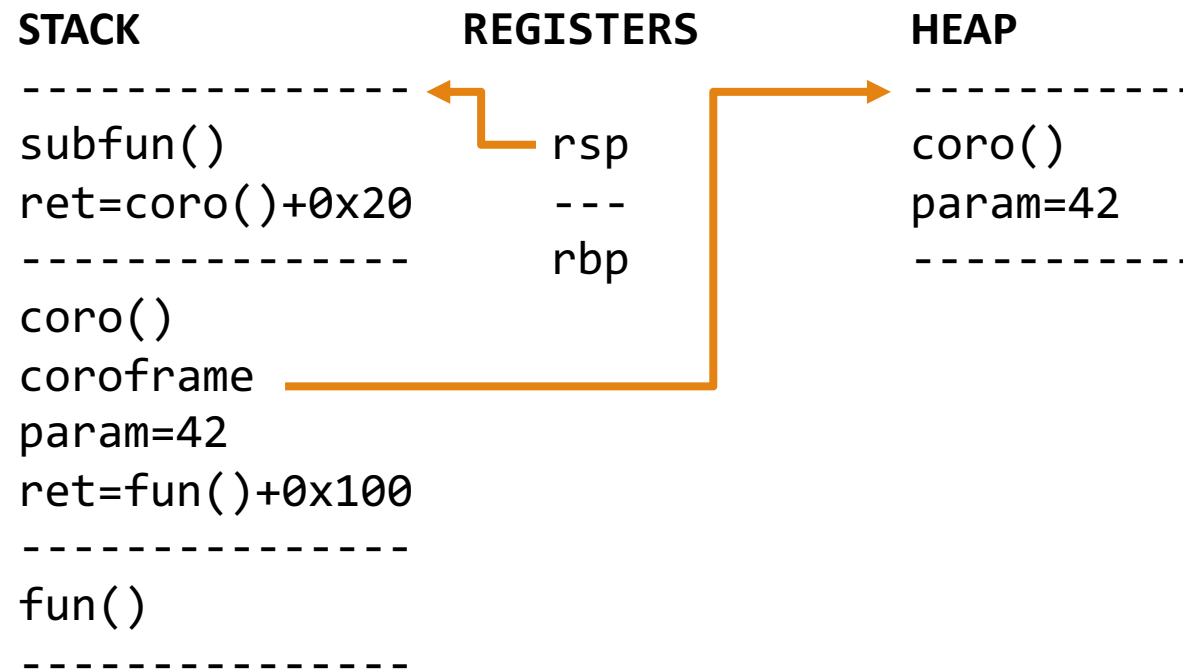
SAMPLE: function **fun()**, calls a coroutine **coro(int param)**

3. coroutine frame allocated



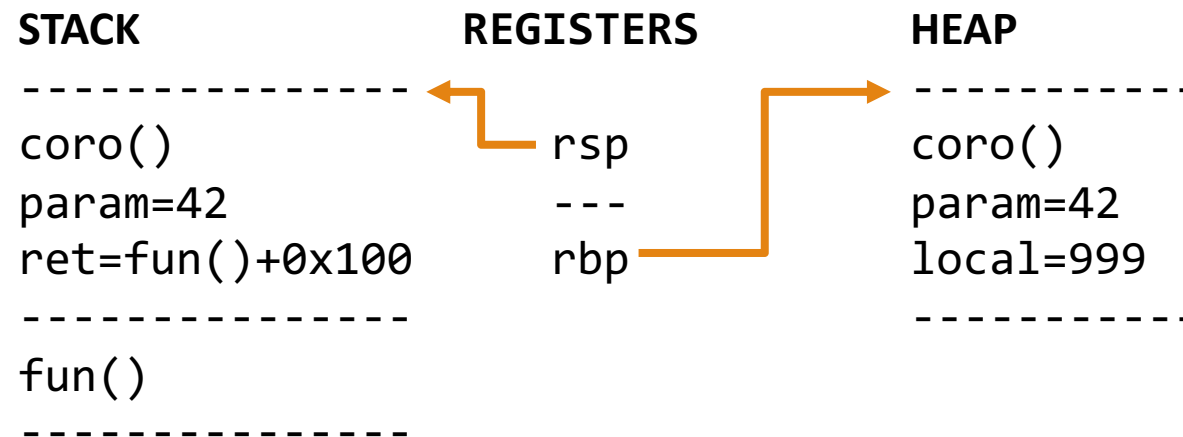
SAMPLE: function **fun()**, calls a coroutine **coro(int param)**

4. Coroutine calls function **subfun()**



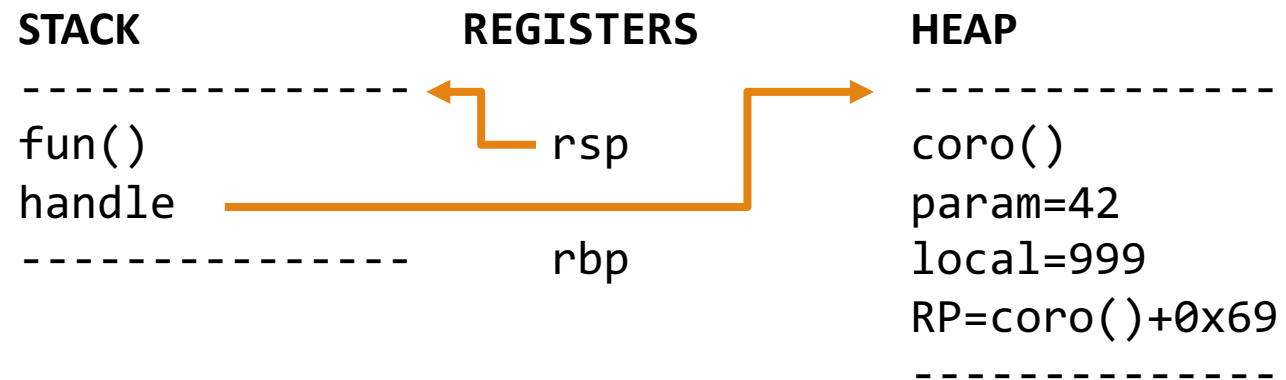
SAMPLE: function **fun()**, calls a coroutine **coro(int param)**

5. **subfun()** returned 999



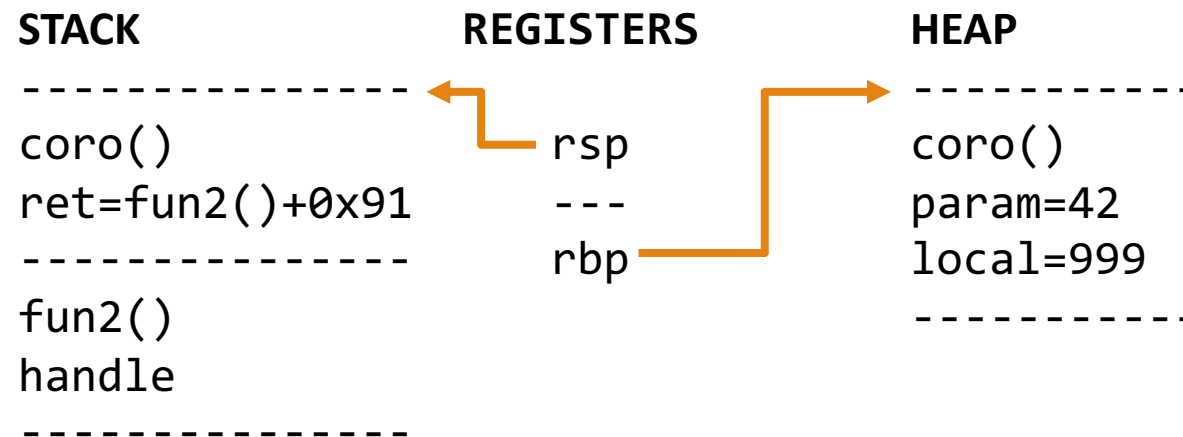
SAMPLE: function **fun()**, calls a coroutine **coro(int param)**


5. **coro()** hits suspension point



SAMPLE: function **fun()**, calls a coroutine **coro(int param)**

5. Some time later... **coro()** gets resumed using handle from different function **fun2()**
(potentially different thread!)





New keywords

Suspending:

- `co_await`
- `co_yield`

Returning:

- `co_return`

Resuming:

- will get back to this...

How to create a Coroutine?

Is this a coroutine?

```
std::future<int> do_some_magic();
```

It depends...

Being a coroutine:

- It's an implementation detail
- Function must use (`co_await`, `co_yield` or `co_return`)
- Has no reflection in function signature/declaration

How to create a Coroutine?

```
std::future<int> do_some_magic()
{
    return std::async([]
    {
        return 42;
    });
}
```

```
std::future<int> do_some_magic()
{
    int result = co_await std::async([]
    {
        return 42;
    });

    co_return result;
}
```

```
void main()
{
    auto ongoing_magic = do_some_magic();
    int magic_number = ongoing_magic.get();
}
```

Cool, but why?

You can create:

- Asynchronous tasks
- Generators
- State machines
- Lock-free barriers
- ...

Simplify writing asynchronous code!

Usig <ppltasks.h>

```
void PickImageClick(Platform::Object^ sender,
Windows::UI::Xaml::RoutedEventArgs^ e)
{
    auto picker = ref new FileOpenPicker();
    picker->FileTypeFilter->Append(L".jpg");
    picker->SuggestedStartLocation = PickerLocationId::PicturesLibrary;

    create_task(picker->PickSingleFileAsync()).then([this]
(Windows::Storage::StorageFile^ file)
    {
        if (nullptr == file)
            return;

        create_task(file->OpenReadAsync()).then([this]
(Windows::Storage::Streams::IRandomAccessStreamWithContentType^ stream)
        {
            auto bitmap = ref new BitmapImage();
            bitmap->SetSource(stream);
            theImage->Source = bitmap;
            OutputDebugString(L"1. End of OpenReadAsync lambda.\r\n");
        });
        OutputDebugString(L"2. End of PickSingleFileAysnc lambda.\r\n");
    });

    OutputDebugString(L"3. End of function.\r\n");
}
```

Using C++ coroutines

```
task<void> PickAnImage()
{
    auto picker = ref new FileOpenPicker();
    picker->FileTypeFilter->Append(L".jpg");
    picker->SuggestedStartLocation = PickerLocationId::PicturesLibrary;

    auto file = co_await picker->PickSingleFileAsync();
    if (nullptr == file)
        return;

    auto stream = co_await file->OpenReadAsync();

    auto bitmap = ref new BitmapImage();
    bitmap->SetSource(stream);
    theImage->Source = bitmap;
    OutputDebugString(L"1. End of function.\r\n");
}
```

Core concepts

Standard introduces new contracts (interfaces) for:

- **Awaitable type**

- Type which could be `co_awaited`

- **Promise type**

- specifies methods for customizing coroutine behavior:
 - what happens when coroutine is called, when returns, etc.
 - customize behavior of `co_await` and `co_yield`
 - communication interface between coroutine and it's caller


Type can be both **Promise** and **Awaitable** at the same time

Awaitable

```
struct awaitable_type
{
    bool await_ready();
    void await_suspend(coroutine_handle<>);
    auto await_resume(); // becomes a return value from co_await expression
};
```

co_await magic

```
auto spell = co_await magic_expr
```



```
auto&& awaitable = magic_expr;  
  
if (!awaitable.await_ready())  
{  
    <suspend coroutine>  
  
    awaitable.await_suspend(coroutine-handle);  
  
    <return-to-caller>  
    <resumption-point>  
}  
  
return awaitable.await_resume();
```

How to resume a coroutine?

Using coroutine handle

```
struct answer_to_life
{
    bool await_ready()
    {
        return false;
    }

    void await_suspend(coroutine_handle<> h) noexcept
    {
        std::thread([h]
        {
            std::this_thread::sleep_for(5s);
            h.resume();
        }).detach();
    }

    auto await_resume()
    {
        return 42;
    }
};
```

```
std::future<int> get_the_answer ()
{
    auto result = co_await answer_to_life(); // 42
    co_return result;
}
```

What's inside of coroutine_handle<>?

```
// Type-erased coroutine handle. Can refer to any kind of coroutine.
// Doesn't allow access to the promise object.
template<>
struct coroutine_handle<void>
{
    // Constructs to the null handle.
    constexpr coroutine_handle();

    // Convert to/from a void* for passing into C-style interop functions.
    constexpr void* address() const noexcept;
    static constexpr coroutine_handle from_address(void* addr);

    // Query if the handle is non-null.
    constexpr explicit operator bool() const noexcept;

    // Query if the coroutine is suspended at the final_suspend point.
    // Undefined behaviour if coroutine is not currently suspended.
    bool done() const;

    // Resume/Destroy the suspended coroutine
    void resume();
    void destroy();
};
```

```
// Coroutine handle for coroutines with a known promise type.
// Template argument must exactly match coroutine's promise type.
template<typename Promise>
struct coroutine_handle : coroutine_handle<>
{
    // Access to the coroutine's promise object.
    Promise& promise() const;

    // You can reconstruct the coroutine handle from the promise object.
    static coroutine_handle from_promise(Promise& promise);
};
```

The Promise

When you write a coroutine that has a body <body-statements>, it's transformed to something like...

```
{
    co_await promise.initial_suspend();
    try
    {
        <body-statements>

        promise.return_value(42); // assuming co_return 42;
        goto FinalSuspend;
    }
    catch (...)
    {
        promise.unhandled_exception();
    }
FinalSuspend:
    co_await promise.final_suspend();
}
```

Putting things together

Let's create a custom coroutine, which can be used like thins:

```
my_custom_coroutine get_the_answer()
{
    auto answer = co_await answer_to_life();
    co_return answer;
}

int main()
{
    auto promise = get_the_answer();
    auto answer = promise.get_result();
    return 0;
}
```

```

using namespace std;
using namespace std::experimental;

struct my_custom_coroutine
{
    struct promise_type
    {
        int _result = 0;

        my_custom_coroutine get_return_object()
        {
            return my_custom_coroutine(coroutine_handle<promise_type>::from_promise(*this));
        }

        auto initial_suspend() { return suspend_never{}; }
        auto final_suspend()   { return suspend_always{}; }

        void return_value(int val)
        {
            _result = val;
        }

        void return_void() {}
    };
    // ...
};

```



Pick one

```

struct my_custom_coroutine
{
    // ...

    coroutine_handle<promise_type>
    my_custom_coroutine(coroutine_handle<promise_type> h)
        : _handle(handle)
    {
    }

    ~my_custom_coroutine()
    {
        if (_handle)
        {
            _handle.destroy();
        }
    }

    int get_result()
    {
        return _handle.promise()._result;
    }
};

```

```

my_custom_coroutine get_the_answer()
{
    auto answer = co_await answer_to_life();
    co_return answer;
}

struct answer_to_life
{
    bool await_ready()
    {
        return false;
    }

    void await_suspend(coroutine_handle<> h) noexcept
    {
        std::thread([h]
        {
            std::this_thread::sleep_for(5s);
            h.resume();
        }).detach();
    }

    auto await_resume()
    {
        return 42;
    }
};

```

```

r();
sult();
for life is: " << result;

```

for life is: 0

```

r();
sult();
for life is: " << result;

```

for life is: 42

Yielding

Can be used for providing multiple values to the caller, not just one!

`co_yield` is just an abstraction over `co_await`

```
generator<int> produce_int(int start, int end)
{
    for(int val = start; val <= end; ++val)
    {
        co_yield val;
    }
}
```

```
int main()
{
    for (auto value : produce_int(5, 10))
    {
        cout << value << endl;
    }
}
```



5
6
7
8
9
10

Magic?

```
int main()
{
    for (auto value : produce_int(5, 10))
    {
        cout << value << endl;
    }
}
```



```
int main()
{
    generator<int> ints = produce_int(5, 10);
    for (auto it = ints.begin(); it != ints.end(); ++it)
    {
        cout << *it << endl;
    }
}
```

What happened there?

```
generator_int produce_int(int start, int end)
{
    for(int val = start; val <= end; ++val)
    {
        co_yield val;
    }
}
```



```
generator_int produce_int(int start, int end)
{
    for(int val = start; val <= end; ++val)
    {
        co_await promise.yield_value(val);
    }
}
```

How to write generator coroutine?

```
struct generator_int
{
    struct promise_type
    {
        const int* _value = nullptr;

        generator_int get_return_object() { /* same as previous*/}

        auto initial_suspend() { return suspend_always{}; }
        auto final_suspend() { return suspend_always{}; }

        auto yield_value(const int& value) ←
        {
            _value = value;
            return suspend_always{};
        }
    };
};

struct iterator : std::iterator<input_iterator_tag, int>
{
    iterator& operator++(); // if not done, resume coroutine
    const int& operator*(); // return current value
};

iterator begin(); // first resumption, initialize iterator with coroutine handle
iterator end();   // return empty iterator
};
```

`void return_value(int val);`
`void return_void();`
`auto yield_value(const int & value);`

C++ Coroutine design principles

Scalable

to billions of concurrent coroutines

Efficient

suspend/resume comparable in cost to function call

Open-ended

library designers can develop coroutine libs exposing high-level semantics

Seamless interaction

with existing facilities with no overhead (e.g. C-style APIs)

Usable

in environments where exceptions are not available or forbidden

Current status

Part of C++20 draft (from Feb 2019)

Supported in:

- Visual Studio 2015 SP2+ (/await)
- Clang 5+ (-fcoroutines-ts -stdlib=libc++)
- GCC – not supported, in progress

Standard library:

- Nothing in standard at this point
- Visual Studio extends `std::future<>` and brings `std::experimental::generator<>`
- [WIP] provide standard implementations (e.g. **CppCoro** library)

Wrapping up

- Coroutine additionally can: **Suspend, Resume** or **Destroy**
- Activation frame is stored on **heap**
- To create coroutine use: **co_await, co_yield** or **co_return** in function implementation
- You need **coroutine handle** to resume it
- To create an expression that could be co_awaited: implement **Awaitable** type contract
- To create custom coroutine type: implement struct with **Promise** type contract
- Use **co_yield** to return multiple values from a single coroutine
- “Just” language feature, use e.g. **cppcoro** to get high-level primitives, such as **generators** or **tasks**
- **Try it at home!**

Useful links

C++ Extensions for Coroutines

<http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2018/n4775.pdf>

Core Coroutines

<http://open-std.org/JTC1/SC22/WG21/docs/papers/2018/p1063r1.pdf>

<https://github.com/lewissbaker/cppcoro>

<https://lewissbaker.github.io/>

“Introduction to C++ Coroutines”:

<https://youtu.be/ZTqHjIm86Bw>

“Coroutine TS a new way of thinking”:

<https://youtu.be/pc-MDA1IXqk>

“Nano-coroutines to the Rescue!”:

<https://youtu.be/j9tIJAqMV7U>

`co_await` questions();

QUIZ TIME!